

Phu Nguyen

The ML workflow felt overwhelming at first, but now I see how each step builds on the last. Splitting the data before training is what makes the evaluation honest. Supervised and unsupervised learning clicked when I realized it's about whether you have the answer during training. I also learned that one accuracy number isn't the whole story. Logistic Regression beat the Decision Tree by only a few wines, and the test set was small. Next time I want to use more features and scale the data.